

Radar-to-Candidate Pipeline TO BE 2.17.4.6.2

Status: Slices 2.17.4.6.2, 2.17.4.7, and 2.17.4.7.0.2 define the implemented retrieval, extraction, language, and evidence boundary. Slice 2.17.4.7.1 adds the approved project-utility scoring and human review lifecycle. Candidate assembly and ranking remain the next target and are not part of scoring.

AS IS sources:

- docs/architecture/RADAR_RUN_PIPELINE_AS_IS.md;
- docs/architecture/UPSTREAM_SEARCH_AND_SIGNAL_ARCHITECTURE.md.

Regenerate PDF:


```
python scripts/generate-draft-run-pipeline-pdf.py `
--source docs/architecture/RADAR_TO_CANDIDATE_PIPELINE_TO_BE_2_17_4_6_2.md `
--output docs/architecture/RADAR_TO_CANDIDATE_PIPELINE_TO_BE_2_17_4_6_2.pdf
```

1. Change Intent

The current RadarRun can plan and execute web search, but its pre-read triage is too dependent on provider order, keyword overlap, exact URL equality, and a first-domain rule. The target pipeline must preserve every discovery decision while reading only the strongest bounded set of sources.

The same architecture must prevent the later signal and candidate stages from repeating the DraftRun context-growth problem. Rich upstream artifacts remain stored for replay and accountability. A provider receives only an operation-specific, budgeted projection with handles back to those artifacts.

2. Target Flow



Nodes through FoundMaterial are implemented by Slice 2.17.4.6.2. Slice 2.17.4.7 implements Signal extraction dossier -> SourceSignal. Slice 2.17.4.7.1 implements SourceSignal -> Signal scoring dossier -> SignalScore and the human review lifecycle. Candidate assembly and ranking remain NOT THIS SLICE and must be implemented by their tracker-backed slices.

3. AS IS to TO BE Mapping

Item	Status	TO BE	Proof
Search campaign planning	UNCHANGED	Deterministic intents, queries, source strategy, and budget skips remain authoritative.	Existing planner tests and trace.
Provider search	CHANGED vs AS IS	Every openWebQuery has a direct current-call budget and final serialized-message proof.	Operation trace, boundary tests, architecture smoke.
Citation normalization	CHANGED vs AS IS	URL, title, and snippet are bounded and normalized without deleting meaningful query parameters.	Normalization tests.
Duplicate handling	CHANGED vs AS IS	Stable duplicate groups retain all query, intent, family, and evidence handles.	Permutation and duplicate-group tests.
Pre-read scoring	CHANGED vs AS IS	Six deterministic dimensions and an explicit quality floor replace keyword-order selection.	Policy tests and readable score trace.
Read allocation	CHANGED vs AS IS	Required-family coverage is allocated first, then quality and bounded diversity.	Allocation and stress tests.
Read failure	CHANGED vs AS IS	Failed reads are failed operations; metadata-only material is not treated as readable.	Integration tests and live trace.
Read-format capability	NEW	The read plan does not spend a slot on an obvious unsupported PDF; unexpected binary responses fail safely into metadata-only evidence.	Reader capability tests and final live trace.
Found material provenance	NEW	discoveryTrace stores resolvable handles without copying rich snippets or full trace objects.	Handle-resolution tests.
Evidence fragments	NEW	Readable materials retain bounded, hashed fragments with offsets before full page text is discarded.	Fragment stability, bounds, and replay tests in Slice 2.17.4.7.
Signal extraction	CHANGED vs AS IS	A backend-owned provider operation receives a bounded extraction dossier, validates exact grounding, and emits zero or more candidate SourceSignals.	Recorded benchmark, provider trace, retry replay, and live proof in Slice 2.17.4.7.
Extraction retry	NEW	A new extraction revision reuses persisted fragments and cannot repeat search or URL reading.	API integration and idempotency proof in Slice 2.17.4.7.
Editorial language context	NEW	BlogProject.language is passed as a bounded project context and remains distinct from source-search languages.	Request-contract, fallback, and trace tests in Slice 2.17.4.7.0.2.
Source-language policy	NEW	Each radar selects editorial-only, editorial-and-English, or unrestricted source eligibility without increasing query count.	Planner allocation, language inspection, triage, and live trace proof in Slice 2.17.4.7.0.2.
Signal localization	CHANGED vs AS IS	Editorial interpretation fields use the project language while source titles and exact evidence quotes remain original. A failed localization emits no mixed-language signal.	Primary/repair/backup language-validation tests and live evidence in Slice 2.17.4.7.0.2.

Item	Status	TO BE	Proof
Project editorial opportunity profile	NEW	A bounded, versioned projection exposes only active project settings and history fingerprints needed for signal utility.	Profile caps, forbidden-field tests, handle resolution, and cross-project benchmark in Slice 2.17.4.7.1.
Signal scoring	CHANGED vs AS IS	Backend-owned batch scoring receives a bounded signal dossier, validates setting/evidence references, and applies a deterministic recommendation policy.	Provider recovery, budget, golden corpus, replay, and live proof in Slice 2.17.4.7.1.
Human signal review	NEW	Utility recommendation and human status remain separate; reversible review events preserve actor, time, reason, revisions, and immutable evidence.	Lifecycle, optimistic-concurrency, API, and authenticated UI tests in Slice 2.17.4.7.1.
Legacy signal integrity	NEW	Old client-evaluated signals are explicitly marked for re-extraction and never displayed as current backend verdicts.	Legacy normalization and UI recovery tests in Slice 2.17.4.7.1.
Candidate assembly and ranking	NOT THIS SLICE	Future assembly receives bounded approved-signal projections.	Slices 2.17.4.8 and 2.17.4.8.1.
Cross-run search memory	NOT THIS SLICE	Reuse of discovered results is owned by a separate durable memory policy.	Slice 2.17.4.6.6.

4. Search Result Triage Contract

The deterministic triage chain is:

`rawResults -> normalized candidates -> duplicate groups -> dimension scores -> read plan -> read outcomes.`

Each raw result receives exactly one terminal decision:

- selected;
- rejected;
- duplicate;
- invalid;
- deferredByBudget.

No result may disappear because of list slicing, duplicate replacement, provider order, or an exception.

4.1 Bounded candidate projection

The triage projection contains only:

- URL, capped at 2048 characters;
- title, capped at 300 characters;
- snippet, capped at 1200 characters;
- query, intent, family, and evidence handles;
- deterministic score dimensions and short reason codes;
- diagnostic explanation, capped at 320 characters.

Full page bodies, complete search responses, source ledgers, previous operation envelopes, and nested budgets are forbidden.

4.2 Normalization and duplicates

Canonical URL normalization removes only known tracking parameters: `utm_*`, `gclid`, `fbclid`, and `ref`. Other query parameters remain because they may identify different documents or views.

Duplicate groups are stable under input permutation. They may be formed by canonical URL, tracking variants, or normalized title/snippet fingerprints. The representative is chosen by deterministic quality and lexical tie-break rules, never by provider position. A group preserves the union of all discovery handles.

4.3 Quality dimensions

Each representative receives scores from 0 to 100 for:

- relevance;
- evidence fit;
- project fit;
- source-quality signals;
- novelty;
- noise risk.

The ordering score is:

$0.30 \text{ relevance} + 0.20 \text{ evidence fit} + 0.20 \text{ project fit} + 0.15 \text{ source quality} + 0.15 \text{ novelty} - \text{noise penalty up to } 30$.

The quality floor is 45. Unknown domains are neutral. Vendor sources are not rejected only for being vendors; pricing and generic promotional noise are penalized through observable content signals.

4.4 Coverage-aware reading

The read allocator first gives every executable required family the best available representative above the quality floor. Remaining capacity is filled by score. Within a score difference of 10, a new evidence type and then a new domain are preferred. Diversity never promotes a candidate below the quality floor.

Read caps remain 1/2/4 for `smoke/standard/full`. Any required direction that could not receive a read is listed in `readCoverageGaps` with a stable reason.

5. Read Outcome and Material Contract

A successful URL read creates a normal readable `FoundMaterial`. A failed read:

- produces a failed URL-read operation;
- preserves the discovery metadata in a `metadataOnly` material;
- records a structured read outcome and warning;
- does not count as a successful readable material.

If every produced material is metadata-only, `RadarRun` status is no better than `partial`.

FoundMaterial.discoveryTrace stores IDs and handles only: raw-result IDs, query IDs, intent IDs, families, evidence types, duplicate-group ID, decision reason, and read outcome. It does not copy snippets, page bodies, or the full searchTriage report.

5.1 Evidence fragment persistence

Before normalized page text is discarded, the reader derives bounded contentFragments. Each fragment has a stable ID, ordinal, text, normalized-text offsets, hash, and semantic kind. Fragments are the smallest persisted evidence unit that a signal may cite. The complete page body remains forbidden in downstream provider input.

A legacy material without fragments may expose its bounded summary as one synthetic fragment. This path is marked legacy-summary-only, has DEGRADED readiness, and cannot produce a trusted high-confidence signal. metadataOnly, unreadable, and empty materials are never sent to the extraction provider.

6. Provider Input Budget Boundary

openWebQuery is the first upstream provider-heavy operation governed by a direct budget contract.

Per call:

Measure	Limit
Query text	1000 characters
Provider input	1500 characters
Serialized messages	4000 characters
Approximate input tokens	1000

Per RadarRun:

Mode	Input characters	Approximate input tokens	Max results per query
smoke	4000	1000	3
standard	12000	3000	5
full	20000	5000	8

The effective result count cannot exceed OPENROUTER_WEB_SEARCH_MAX_RESULTS. Before the provider call, the direct input gate checks the current query and run totals. After message construction, a provider-neutral message guard measures the actual serialized messages. An over-limit operation is not sent and records provider-input-over-budget plus a structured incident.

The operation trace contains the actual providerInput, payloadBudget, inputStats, payloadStats, messageCharCount, model selection, and provider usage when returned. Nested metadata from an older artifact never counts as proof.

Provider-reported prompt tokens can exceed the local message estimate because the OpenRouter web-search tool adds its own retrieval context. This value is preserved as usage/cost telemetry, but it is not confused with the Glavred-controlled provider input. Production limits for provider-owned search cost and reuse belong to Slice 2.17.4.6.6.

6.1 Editorial and source language boundary

The project-owned editorial language and the radar-owned source-language policy are separate inputs. The API passes only projectId and editorialLanguage; a provider never receives the complete portfolio project. A legacy request without this bounded context uses a trace-visible fallback rather than silently claiming canonical project metadata.

Each radar resolves one policy:

- `editorialOnly`: all query families use the editorial language and a confidently detected different source language is ineligible;
- `editorialAndEnglish`: broad discovery, limitations, and freshness use the editorial language, while case, benchmark, and OSS families use English; when the editorial language is English, all families use English;
- `any`: the same bounded query allocation is used, but source language does not restrict eligibility.

The allocation never duplicates a query family and never increases `maxExternalQueries`. A budget that cannot execute every planned query language records `languageCoverageGaps`. A deterministic inspector classifies bounded search metadata and retained fragments as a primary language, confidence, and mixed/unknown state. Only a high-confidence disallowed language is rejected; mixed and unknown content continue with a warning so that detection uncertainty cannot silently remove evidence.

7. Signal Extraction Contract

Signal extraction is a backend upstream operation owned separately from project utility scoring. It answers what evidence-backed fact, change, tension, case, data point, practice, failure mode, observation, question, or recurring pattern exists in the material. It does not choose a topic, fabula, audience, value, goal, platform, or publication channel.

The rich input is the persisted set of readable `FoundMaterial` records and their fragments. `SignalExtractionContextFactory` produces a bounded radar context from scope, active rules, source intent, evidence types, filter references, and the language context. `SignalExtractionDossierFactory` then retains only:

- material IDs, source metadata, and selected bounded fragments;
- radar rule/filter references needed to understand the search scope;
- the extraction taxonomy and required output contract;
- handles back to persisted material and fragment artifacts.

Full workspace snapshots, complete pages, topics, fabulas, plans, publication history, previous envelopes, and nested budget artifacts are `neverSendToProvider`.

7.1 Terminal material decisions

Every inspected material receives exactly one decision: `signalProducing`, `insufficient`, `duplicate`, `corroborating`, `contradiction`, `noise`, or `extractionFailed`.

One material may produce zero, one, or several signals. Several materials may support or contradict a canonical signal. Signal count is never a target. Every accepted signal must resolve to at least one exact retained fragment.

7.2 Grounding and recovery

`SignalGroundingPolicy` rejects unknown handles, quotations absent from retained fragments, changed numbers or dates, unsupported certainty, and invented actors, mechanisms, outcomes, or limitations. A malformed primary response is repaired by the same model using only structured errors, then attempted with the backup model. If all provider paths fail, the terminal fallback emits no substantive signals and marks the affected materials `extractionFailed`.

Retrieval and extraction statuses are independent. Successful search and reading remain successful when extraction is partial, failed, or not run. A manual retry creates a new report revision from persisted fragments and performs no search or URL read operations.

7.3 Extraction budget boundary

Mode	Materials	Fragments per material	Fragment characters	Provider input	Serialized messages	Approximate input tokens	Max output tokens
smoke	1	3	700	6000	9000	2250	1200
standard	2	4	900	12000	16000	4000	2200
full	4	5	900	24000	30000	7500	3500

Primary, repair, and backup attempts each pass a direct current-call input gate and a final serialized-message guard. Repair context is at most 1200 characters and is included in both checks. An over-budget attempt never calls the provider. Actual OpenRouter usage is stored when supplied; missing provider usage remains explicitly unknown.

7.4 Editorial localization and original evidence

title, summary, uncertainty, mechanism, outcome, and limitations are editorial interpretation fields and must use `editorialLanguage`. Source titles and exact evidence quotations remain in the original source language and are never translated by this operation.

A deterministic language policy validates non-empty editorial prose while ignoring URLs, IDs, numbers, model names, and short technical abbreviations. A language violation is a structured payload error: primary is followed by same-model repair and then backup under the existing attempt and budget policy. If all attempts violate the contract, no `SourceSignal` is emitted; the material remains in the extraction report as `extractionFailed` with `editorial-language-not-satisfied` and can be retried from persisted fragments.

Accepted signals expose editorial language, detected source language, localization status, and reason codes. Evidence IDs include signal, material, fragment, and quote identity so multiple exact quotes from one fragment remain stable and unique.

8. Signal Utility Scoring and Review Contract

Signal scoring answers whether an extracted, evidence-backed signal is useful for a specific editorial project. It does not change source evidence, approve the signal, select a topic or fabula, create a `PostCandidate`, or reserve a plan slot.

The runtime chain is:

```
SourceSignal -> ProjectEditorialOpportunityProfile -> SignalUtilityDossier -> provider
evaluation -> deterministic decision policy -> SignalUtilityReport -> human review lifecycle.
```

8.1 Bounded project profile and signal dossier

`ProjectEditorialOpportunityProfileFactory` derives a traceable projection from the canonical workspace. Standard mode retains at most 24 active editorial rules, 16 confirmed author-position assertions, 12 active topics, and 20 fingerprints from previous signals, candidates, and publications. Every retained setting has a stable ID that resolves to the workspace. Disabled settings are not silently promoted.

`SignalUtilityDossierFactory` adds at most six signals per standard batch and four evidence references per signal. The provider sees only editorial language, a short project profile, active bounded settings, filter contracts, signal interpretation fields, and resolvable evidence handles. The full workspace, fabulas, content plan, complete publications, old traces, and provider envelopes are forbidden.

Each utility dimension has one owner and one result status: `matched`, `partial`, `notProven`, or `conflict`. The supported dimensions cover evidence strength, factual specificity, source credibility, mechanism, observable outcome, actionability, topic/author/audience/positioning/goal fit, novelty, productive tension, freshness, duplication/banality, promotional noise, and prohibited content. Legacy dimensions are mapped to these owners only for compatibility.

8.2 Deterministic recommendation policy

The provider performs semantic comparison and returns dimension evidence. It cannot approve, reject, archive, or correct a signal. The backend validates signal IDs, setting IDs, material/fragment evidence handles, dimension ownership, and required reasons before computing one terminal recommendation:

- `recommended`: evidence is sufficient and all blocking criteria are satisfied;
- `reviewWithCaution`: no proven blocker exists, but evidence, source independence, optional fit, or uncertainty needs editorial attention;
- `notRecommended`: an active blocking criterion has a proven conflict;
- `inconclusive`: provider recovery failed or references are insufficient for an honest decision.

Filter modes determine importance: `mustMatch` and `mustNotMatch` are blocking, `shouldMatch` is weighted, and `seekTension` is diagnostic. `notProven` is never a blocking rejection by itself. No aggregate percentage may hide blockers, missing proof, or uncertainty.

8.3 Provider recovery and budget boundary

Scoring uses one bounded provider batch per revision, followed when necessary by a same-model repair and backup attempt. Repair context contains only structured validation errors and is capped at 1200 characters. Exhausted recovery yields `inconclusive`; retrieval and extraction remain successful.

Mode	Signals	Provider input	Serialized messages	Approximate input tokens	Max output tokens
smoke	1	6000	9000	2250	1200
standard	6	16000	22000	5500	3200
full	12	28000	36000	9000	5000

Every primary, repair, and backup attempt has direct current-call input proof and a final serialized-message check. An over-budget or blocked dossier never calls the provider. Trace stores the exact bounded provider input, profile/dossier metadata, retained/trimmed/suppressed counts, message size, actual provider usage when available, validation incidents, unresolved-reference counts, and policy result.

The provider response uses compact aliases declared by the dossier: `signalKey`, `criterionKey`, and `evidenceKeys`. The backend resolves them to persisted signal, setting, material, and fragment handles before accepting the response. Aliases reduce output growth but never weaken identity validation or allow positional inference.

8.4 Human review lifecycle

Human status is independent of utility recommendation. The initial status is `candidate` (На проверке). Allowed transitions are `candidate` or `corrected` to `approved/rejected/archived`, `approved` or `rejected` back to `candidate`, and `archived` back to `candidate`. Correction changes only editorial title, summary, and author comment. Rejection, archive, and correction require a reason.

Each immutable event records ID, authenticated actor, timestamp, action, from/to status, reason, changed fields, and expected review revision. Evidence, exact quotes, mechanism, outcome, limitations, and provenance cannot be edited. A correction marks the utility report stale and triggers a bounded rescore; failed rescore does not undo the saved correction.

8.5 Legacy and presentation integrity

Loading a legacy signal never triggers a hidden provider call. A signal without the language/scoring contract is marked `needsReExtraction`; any old client keyword evaluation is retained only as explicitly historical diagnostics. It cannot be rendered as the current recommendation or human decision.

The Signals UI keeps source, evidence, extraction semantics, utility dimensions, and review history visible while editing. Only editorial title, summary, and author comment are editable. Source titles, quotations, and radar names wrap inside their owners and remain fully reachable at supported desktop and mobile widths. Radar rows present full title/status, description, and operational metadata as separate layout levels rather than squeezing them into one line.

8.6 Explainable criteria and signal relationships

`SignalUtilityReport` version 2 no longer presents one flat list of dimensions as if all checks had the same origin. It separates:

- one criterion for every enabled radar filter;
- applicable author, audience, positioning, project-goal, topic, and prohibition

criteria from the bounded project profile;

- type-aware system checks for evidence grounding, result support, source posture, freshness, and relationship integrity.

Every criterion snapshots its human-readable title, statement, mode, verdict, effect, explanation, setting handles, evidence handles, and uncertainty. Retained settings that are not applicable to signal utility are listed explicitly as `notApplicable` or `notThisStage`. Style, channel, fabula, and plan mechanics remain downstream.

The signal type owns which semantic fields and system checks are applicable. A non-empty mechanism or outcome is not proof by itself. Result support distinguishes observed, reported, `capabilityOnly`, expected, missing, and `notApplicable`. Absence of a known vendor marker never proves source independence; source posture is reported as independent, corroborated, `firstParty`, vendor, or unknown.

Internal material, fragment, and setting IDs remain trace handles. The user-facing report resolves them to the criterion text, source title/domain, exact quotation, source URL, and trace link. An unresolved handle makes the affected result inconclusive; raw IDs are not presented as editorial evidence.

Signal relationship analysis is independent from project utility. A bounded hybrid policy classifies candidate pairs as `exactDuplicate`, `sameClaim`, `relatedSameSource`, corroborates, contradicts, distinct, or inconclusive. Exact duplicates and same-claim aliases share one canonical presentation while all source records and evidence remain persisted. Related claims from the same material remain separate and visibly linked. Missing relationship proof is `notChecked` or inconclusive, never an invented low duplicate risk.

Relationship candidates reuse the existing scoring provider attempt and budget. They do not add an unbounded provider call or raise the current provider-input and serialized-message limits.

9. Future Provider Context Rule

Every future upstream provider-heavy stage must declare before implementation:

1. a typed rich input artifact;
2. an operation-specific dossier/input owner;
3. `mustHave`, `shouldHave`, `diagnosticOnly`, and `neverSendToProvider` fields;
4. a direct current-call budget profile;
5. a final serialized-message guard;
6. handles back to persisted artifacts;
7. a stress test proving bounded growth;
8. a trace-safe outcome and fallback/incident policy.

Architecture smoke rejects a new operation that does not satisfy this inventory or carry an explicit tracker-backed debt exception.

10. Trace Contract

RadarRun keeps existing fields and adds `searchTriage`:

- policy version;
- normalized candidates and dimension scores;
- duplicate groups;
- read plan and terminal decisions;
- required-family coverage and gaps;
- read outcomes;
- terminal decision counts.

Existing `rawResults`, `selectedForRead`, and `rejectedBeforeRead` remain compatible. Old runs without `searchTriage` remain readable in the UI.

Slice 2.17.4.7.0.2 additionally records the bounded language context, per-intent and per-query language, source-language inspection, eligibility reasons, query-language coverage gaps, and signal localization status. Compatibility `searchPlan.language` continues to mean editorial language.

Slice 2.17.4.7 additionally stores `run.signalExtraction`, `signalExtractionReport`, and `sourceSignals`. The report includes revision history, material decisions, grounding violations, duplicate/corroboration/contradiction links, provider attempts, direct budgets, final message sizes, usage, and suppressed fields.

`FoundMaterial.contentFragments` are persisted evidence artifacts, not copied trace prose.

Slice 2.17.4.7.1 additionally stores `run.signalScoring`, `signalScoringReport`, source-signal utility reports, review revisions, and immutable review events. Manual scoring reuses stored signals/materials and performs no search, URL read, or extraction. Legacy heuristic results remain explicitly historical.

11. Success Criteria

- Every raw result has one terminal decision.
- Selection and duplicate representatives are invariant under provider result order.
- No duplicate group schedules more than one URL read.
- Required-family coverage is maximized inside the read cap and gaps are explicit.
- Known generic-news and pricing noise does not displace a suitable alternative.

- Failed reads remain failed and metadata-only results are not counted as readable.
- One hundred raw results cannot grow provider calls, provider messages, or read count beyond the active profile.
- openWebQuery has direct input and final-message budget proof.
- Retrieval may create candidate SourceSignal artifacts only through the extraction owner. It still creates no PostCandidate, plan slot, editorial work item, or DraftRun.
- Every inspected material has one terminal extraction decision and every accepted signal resolves to exact material and fragment handles.
- Manual extraction retry is idempotent and never repeats search or URL reading.
- Unsupported certainty, altered numbers/dates, and unresolved evidence handles are zero in the accepted benchmark and live proof.
- The canonical editorial language comes from BlogProject.language; source search policy and source language are separate trace fields.
- Language policy changes actual bounded query allocation without adding provider calls, and every skipped query language has an explicit coverage gap.
- Accepted editorial fields use the project language while exact source titles and evidence quotes remain original.
- A terminal localization failure creates no mixed-language signal, and every displayed evidence item resolves to a safe source URL plus material/fragment trace.
- Every scored signal has one terminal recommendation or explicit inconclusive; every dimension has a specific reason and resolvable setting/evidence references.
- A signal is notRecommended only for a proven conflict with an active blocking criterion. Missing keyword overlap and notProven never become automatic rejection.
- Vendor-only evidence, credible off-project evidence, weak evidence, productive tension, and prohibited content remain distinguishable outcomes.
- Utility recommendations never change human review status. Review transitions are authenticated, reversible, revision-checked, and preserve evidence exactly.
- Provider scoring remains within direct input and final-message caps, and manual rescore performs no retrieval or extraction work.
- Legacy client evaluations are not displayed as current backend recommendations.
- Signals and radar rows preserve complete readable information without page-level horizontal overflow at the supported five viewport widths.
- Recorded and comparable live proof show no quality regression relative to the pre-change industrial-AI baseline.

12. Implementation Status

- Slice 2.17.4.6.2: triage v2, selective reading, read-outcome trace, upstream budget boundary, UI trace, recorded/live proof. IMPLEMENTED.
- Slice 2.17.4.7: evidence fragments, bounded provider extraction, grounding,

material decisions, retry, UI trace, recorded/live proof. IMPLEMENTED.

- Slice 2.17.4.7.0.2: project/editorial language handoff, radar source-language

policy, query allocation, source-language eligibility, signal localization, and evidence presentation. IMPLEMENTED; accepted live proof is recorded below.

- Slice 2.17.4.7.1: bounded project profile, utility dossier, backend-owned scoring,

deterministic recommendation, human review lifecycle, legacy/UI integrity repair. IMPLEMENTED; accepted runtime proof is recorded in docs/evidence/radar-runs/2.17.4.7.1/.

- Candidate assembly, candidate ranking, and cross-run search memory: NOT THIS SLICE.

13. Implementation Proof

The final live proof on 2026-07-13 returned 52 raw results and exactly 52 terminal decisions: 2 selected, 7 rejected, 17 duplicate, and 26 deferred by budget. It built 35 stable duplicate groups and produced two readable materials from two domains with no read warnings, accepted noise, or metadata-only output.

All three openWebQuery operations were directly budgeted. Serialized messages used 1,185 characters in total, local approximate input was 297 tokens, and no budget incident was recorded. The benchmark remained warning only because the existing three-query campaign budget skipped required limitationCritique; required coverage was not lost by triage.

Trace-safe pre/post evidence is committed in docs/evidence/radar-runs/2.17.4.6.2/COMPARISON.md and comparison.json.

The accepted extraction proof on 2026-07-14 is radar-run-ai-pattern-radar-industrial-cases-8. It returned 60 raw results, read two materials from two domains, and created three grounded signal candidates. Both materials received a terminal extraction decision; unresolved evidence handles and downstream artifacts were zero.

The initial extraction used 12,496 serialized characters against the 16,000 standard cap and 3,743 provider-reported tokens. A forced retry did not add search or URL-read operations. Its first response was rejected for an ungrounded number, same-model repair was accepted, and both attempts remained below the message cap. The live retrieval benchmark stayed warning only because the existing three-query budget skipped limitationCritique.

Trace-safe evidence is committed in docs/evidence/radar-runs/2.17.4.7/BASELINE.md and LIVE_PROOF.md with JSON peers.

The accepted language-policy proof is RadarRun radar-run-ai-pattern-radar-industrial-cases-3, extraction revision 4. The bounded project context resolved editorialLanguage=ru and sourceLanguagePolicy=editorialAndEnglish. The unchanged three-query budget assigned Russian to broad discovery and English to case and benchmark families; it added no provider call and reported no language coverage gap.

The run retained one Russian and one English readable material. The English source produced three Russian candidate signals while preserving the original source title, exact English quotes, URLs, and resolvable material/fragment handles. The first extraction response was rejected because one quote did not exactly match its source fragment. Same-model repair was accepted after receiving a 304-character structured correction context. The two attempts used 11,972 and 12,337 serialized characters against the unchanged 16,000 cap, with 8,367 actual tokens in total and no budget incident. The rejected grounding incident remains visible in attempt diagnostics, while the accepted signals have no unresolved or unsupported evidence handles.

The authorized UI proof also verified backend persistence, Ha npoверке, the unscored utility state, original source navigation, deep trace navigation, and the localized demo radar. Trace-safe evidence is committed in docs/evidence/radar-runs/2.17.4.7.0.2/.

The accepted scoring/explainability proof on 2026-07-22 uses fresh RadarRun radar-run-ai-pattern-radar-industrial-cases-9. It produced three Russian signals; all three matched the industrial radar and received reviewWithCaution. Automatic scoring was accepted on the primary attempt with 17,166 serialized characters against the 22,000 cap and 4,502/1,308 input/output tokens. Setting and evidence references resolved with zero errors. Manual rescore created revision 2 without search, read, or extraction work, and one authenticated review event approved a signal without mutating its evidence.

Retrieval was partial because one search response contained malformed JSON and one OpenReview URL returned HTTP 403. Both incidents remain explicit in operation and benchmark trace; they did not turn successful extraction or scoring into a false clean retrieval result.

Replay radar-run-ai-pattern-radar-industrial-cases-7 classified the digital-advisor result as capabilityOnly, kept source posture unknown, and linked it to the separate automated-risk thesis as relatedSameSource. Product UI hid raw handles, opened source and trace links, and stayed within the page width at 390, 1180, 1440, 1904, and 2048 pixels. Trace-safe evidence is committed in docs/evidence/radar-runs/2.17.4.7.1/.